

```

7974 NodeManager
8168 org.eclipse.equinox.launcher_1.6.400.v20210924-0641.jar
7641 ResourceManager
7482 SecondaryNameNode
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls
ls: '.': No such file or directory
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
Found 19 items
drwxr-xr-x - hduser supergroup          0 2023-05-11 14:25 /CSE
drwxr-xr-x - hduser supergroup          0 2023-05-11 13:58 /FFF
drwxr-xr-x - hduser supergroup          0 2023-05-11 14:31 /LLL
drwxr-xr-x - hduser supergroup          0 2019-10-23 16:07 /STUDENT_INFO
drwxr-xr-x - hduser supergroup          0 2023-05-11 13:54 /aftfah
drwxr-xr-x - hduser supergroup          0 2019-10-23 15:08 /arv
drwxr-xr-x - hduser supergroup          0 2023-05-11 15:00 /first
drwxr-xr-x - hduser supergroup          0 2023-05-04 13:05 /inputbda
drwxr-xr-x - hduser supergroup          0 2023-04-27 11:48 /labShadoop
drwxr-xr-x - hduser supergroup          0 2023-05-08 09:40 /new_folder
drwxr-xr-x - hduser supergroup          0 2022-06-14 10:14 /output
drwxr-xr-x - hduser supergroup          0 2023-05-04 13:15 /outputbda
drwxr-xr-x - hduser supergroup          0 2022-06-14 10:09 /rgs
drwxr-xr-x - hduser supergroup          0 2023-05-12 11:58 /tempinput
-rw-r--r-- 1 hduser supergroup      888190 2023-05-12 12:55 /tempinput1
drwxr-xr-x - hduser supergroup          0 2023-05-12 12:56 /tempout
drwxr-xr-x - hduser supergroup          0 2023-04-27 11:47 /test
drwxrwxr-x - hduser supergroup          0 2019-10-23 15:36 /tmp
drwxr-xr-x - hduser supergroup          0 2019-08-01 16:03 /user
hduser@bmsce-Precision-T1700:~$ hadoop fs -mkdir /rgs
mkdir: '/rgs': File exists
hduser@bmsce-Precision-T1700:~$ hadoop fs -copyFromLocal /home/hduser/Desktop/sample.txt /rgs/test.txt
copyFromLocal: '/rgs/test.txt': File exists
hduser@bmsce-Precision-T1700:~$ hadoop fs -copyFromLocal /home/hduser/Desktop/sample.txt /rgs/test1.txt
hduser@bmsce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/WordCount.jar WCDriver /rgs/test1.txt /home/hduser/Desktop/abc.txt
Exception in thread "main" java.lang.ClassNotFoundException: WCDriver
    at java.net.URLClassLoader.findClass(URLClassLoader.java:382)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:418)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:351)
    at java.lang.Class.forName0(Native Method)
    at java.lang.Class.forName(Class.java:348)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:214)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/WordCount.jar WCDriver /rgs/test1.txt /home/hduser/Desktop/abc.txt
23/05/17 10:07:35 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
23/05/17 10:07:35 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
23/05/17 10:07:35 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker, sessionId= - already initialized
23/05/17 10:07:35 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with Tool
23/05/17 10:07:35 INFO mapred.FileInputFormat: Total input paths to process : 1
23/05/17 10:07:35 INFO mapreduce.JobSubmitter: number of splits:1
23/05/17 10:07:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local713144167_0001
23/05/17 10:07:36 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
23/05/17 10:07:36 INFO mapred.LocalJobRunner: OutputCommitter set in config null
23/05/17 10:07:36 INFO mapreduce.Job: Running job: job_local713144167_0001
23/05/17 10:07:36 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
23/05/17 10:07:36 INFO mapred.LocalJobRunner: Waiting for map tasks
23/05/17 10:07:36 INFO mapred.LocalJobRunner: Starting task: attempt_local713144167_0001_m_000000_0

```

hduser@bmsce-Precision-T1700: ~

```

Map-Reduce Framework
  Map input records=5
  Map output records=84
  Map output bytes=504
  Map output materialized bytes=678
  Input split bytes=88
  Combine input records=0
  Combine output records=0
  Reduce input groups=17
  Reduce shuffle bytes=678
  Reduce input records=84
  Reduce output records=17
  Spilled Records=168
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=2
  CPU time spent (ms)=0
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=466616320
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=89
File Output Format Counters
  Bytes Written=70

```

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /home/hduser/Desktop/abc.txt
```

```
Found 2 items
```

```

-rw-r--r-- 1 hduser supergroup      0 2023-05-17 10:07 /home/hduser/Desktop/abc.txt/_SUCCESS
-rw-r--r-- 1 hduser supergroup    70 2023-05-17 10:07 /home/hduser/Desktop/abc.txt/part-000000

```

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -cat /home/hduser/Desktop/abc.txt/part-000000
```

```

15
2
2
3
1
7
7
1
1
1
12
7
6
3
5
5
6

```

```
hduser@bmsce-Precision-T1700:~$
```



```

May 17 10:16
hduser@bmsce-Precision-T1700: ~
hduser@bmsce-Precision-T1700:~$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [localhost]
hduser@localhost's password:
localhost: namenode running as process 7093. Stop it first.
hduser@localhost's password:
localhost: datanode running as process 7268. Stop it first.
Starting secondary namenodes [0.0.0.0]
hduser@0.0.0.0's password:
0.0.0.0: secondarynamenode running as process 7482. Stop it first.
starting yarn daemons
resourcemanager running as process 7641. Stop it first.
hduser@localhost's password:
localhost: nodemanager running as process 7974. Stop it first.
hduser@bmsce-Precision-T1700:~$ jps
9937 Jps
7268 DataNode
7093 NameNode
7974 NodeManager
8168 org.eclipse.equinox.launcher_1.6.400.v20210924-0641.jar
7641 ResourceManager
7482 SecondaryNameNode
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls
ls: '.': No such file or directory
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
Found 19 items
drwxr-xr-x - hduser supergroup          0 2023-05-11 14:25 /CSE
drwxr-xr-x - hduser supergroup          0 2023-05-11 13:58 /FFF
drwxr-xr-x - hduser supergroup          0 2023-05-11 14:31 /LLL
drwxr-xr-x - hduser supergroup          0 2019-10-23 16:07 /STUDENT_INFO
drwxr-xr-x - hduser supergroup          0 2023-05-11 13:54 /aflfah
drwxr-xr-x - hduser supergroup          0 2019-10-23 15:08 /arv
drwxr-xr-x - hduser supergroup          0 2023-05-11 15:00 /flrst
drwxr-xr-x - hduser supergroup          0 2023-05-04 13:05 /inputbda
drwxr-xr-x - hduser supergroup          0 2023-04-27 11:48 /lab5hadoop
drwxr-xr-x - hduser supergroup          0 2023-05-08 09:40 /new_folder
drwxr-xr-x - hduser supergroup          0 2022-06-14 10:14 /output
drwxr-xr-x - hduser supergroup          0 2023-05-04 13:15 /outputbda
drwxr-xr-x - hduser supergroup          0 2022-06-14 10:09 /rgs
drwxr-xr-x - hduser supergroup          0 2023-05-12 11:58 /tempinput
-rw-r--r-- 1 hduser supergroup      888190 2023-05-12 12:55 /tempinput1
drwxr-xr-x - hduser supergroup          0 2023-05-12 12:56 /tempout
drwxr-xr-x - hduser supergroup          0 2023-04-27 11:47 /test
drwxrwxr-x - hduser supergroup          0 2019-10-23 15:36 /tmp
drwxr-xr-x - hduser supergroup          0 2019-08-01 16:03 /user
hduser@bmsce-Precision-T1700:~$ hadoop fs -mkdir /rgs
mkdir: '/rgs': File exists
hduser@bmsce-Precision-T1700:~$ hadoop fs -copyFromLocal /home/hduser/Desktop/sample.txt /rgs/test.txt
copyFromLocal: '/rgs/test.txt': File exists
hduser@bmsce-Precision-T1700:~$ hadoop fs -copyFromLocal /home/hduser/Desktop/sample.txt /rgs/test1.txt
hduser@bmsce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/WordCount.jar WCDriver /rgs/test1.txt /home/hduser/Desktop/abc.txt
Exception in thread "main" java.lang.ClassNotFoundException: WCDriver
    at java.net.URLClassLoader.findClass(URLClassLoader.java:382)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:418)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:351)
    at java.lang.Class.forName0(Native Method)

```



USN : 1BM20CS195

```
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~$ jadoop fs -copyFromLocal /home/hduser/Desktop/1901 /rgs/test2.txt
jadoop: command not found
hduser@bmsce-Precision-T1700:~$ jadoop fs -copyFromLocal /home/hduser/Desktop/1901 /rgs/test2.txt
hduser@bmsce-Precision-T1700:~$ jadoop jar /home/hduser/Desktop/AverageTemperature.jar AverageDriver /rgs/test2.txt /home/hduser/Desktop/abc1.txt
23/05/17 10:53:57 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
23/05/17 10:53:57 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
23/05/17 10:53:57 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
23/05/17 10:53:57 INFO input.FileInputFormat: Total input paths to process : 1
23/05/17 10:53:57 INFO mapreduce.JobSubmitter: number of splits:1
23/05/17 10:53:57 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1023094052_0001
23/05/17 10:53:57 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
```

```
FILE: Number of read operations=0
FILE: Number of large read operations=0
HDFS: Number of write operations=0
HDFS: Number of bytes read=1776380
HDFS: Number of bytes written=8
HDFS: Number of read operations=13
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
Map-Reduce Framework
  Map input records=6565
  Map output records=6564
  Map output bytes=59076
  Map output materialized bytes=72210
  Input split bytes=101
  Combine input records=0
  Combine output records=0
  Reduce input groups=1
  Reduce shuffle bytes=72210
  Reduce input records=6564
  Reduce output records=1
  Spilled Records=13128
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=61
  CPU time spent (ms)=0
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=999292928
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=888190
File Output Format Counters
  Bytes Written=8
```

```
hduser@bmsce-Precision-T1700:~$ jadoop fs -ls /home/hduser/Desktop/abc1.txt
Found 2 items
-rw-r--r-- 1 hduser supergroup 0 2023-05-17 10:53 /home/hduser/Desktop/abc1.txt/_SUCCESS
-rw-r--r-- 1 hduser supergroup 8 2023-05-17 10:53 /home/hduser/Desktop/abc1.txt/part-r-00000
hduser@bmsce-Precision-T1700:~$ ^C
hduser@bmsce-Precision-T1700:~$ jadoop fs -cat /home/hduser/Desktop/abc1.txt/part-r-00000
1901 46
hduser@bmsce-Precision-T1700:~$
```